

Tree Methods Project

Guan-Yuan Wang

2020/9/3

Tree Methods Project

For this project we will be exploring the use of tree methods to classify schools as Private or Public based off their features.

Let's start by getting the data which is included in the ISLR library, the College data frame.

A data frame with 777 observations on the following 18 variables.

- Private A factor with levels No and Yes indicating private or public university
- Apps Number of applications received
- Accept Number of applications accepted
- Enroll Number of new students enrolled
- Top10perc Pct. new students from top 10% of H.S. class
- Top25perc Pct. new students from top 25% of H.S. class
- F.Undergrad Number of fulltime undergraduates
- P.Undergrad Number of parttime undergraduates
- Outstate Out-of-state tuition
- Room.Board Room and board costs
- Books Estimated book costs
- Personal Estimated personal spending
- PhD Pct. of faculty with Ph.D.'s
- Terminal Pct. of faculty with terminal degree
- S.F.Ratio Student/faculty ratio
- perc.alumni Pct. alumni who donate
- Expend Instructional expenditure per student
- Grad.Rate Graduation rate

Get the Data

Call the ISLR library and check the head of College (a built-in data frame with ISLR, use `data()` to check this.) Then reassign College to a dataframe called `df`

```
library(ISLR)
```

```
## Warning: package 'ISLR' was built under R version 3.6.3
```

```
#data()
head(College)
```

```
##               Private Apps Accept Enroll Top10perc Top25perc
## Abilene Christian University    Yes 1660  1232    721      23      52
## Adelphi University              Yes 2186  1924    512      16      29
## Adrian College                 Yes 1428  1097    336      22      50
## Agnes Scott College             Yes  417   349    137      60      89
## Alaska Pacific University       Yes  193   146     55      16      44
## Albertson College              Yes  587   479    158      38      62
##               F.Undergrad P.Undergrad Outstate Room.Board Books
## Abilene Christian University    2885      537    7440      3300   450
## Adelphi University             2683      1227   12280      6450   750
## Adrian College                 1036        99   11250      3750   400
## Agnes Scott College             510        63   12960      5450   450
## Alaska Pacific University       249      869    7560      4120   800
## Albertson College              678        41   13500      3335   500
##               Personal PhD Terminal S.F.Ratio perc.alumni Expend
## Abilene Christian University    2200  70      78     18.1      12   7041
## Adelphi University             1500  29      30     12.2      16  10527
## Adrian College                 1165  53      66     12.9      30   8735
## Agnes Scott College             875  92      97      7.7      37  19016
## Alaska Pacific University       1500  76      72     11.9       2  10922
## Albertson College              675  67      73      9.4      11   9727
##               Grad.Rate
## Abilene Christian University    60
## Adelphi University             56
## Adrian College                 54
## Agnes Scott College            59
## Alaska Pacific University       15
## Albertson College              55
```

```
df <- College
```

EDA

Let's explore the data!

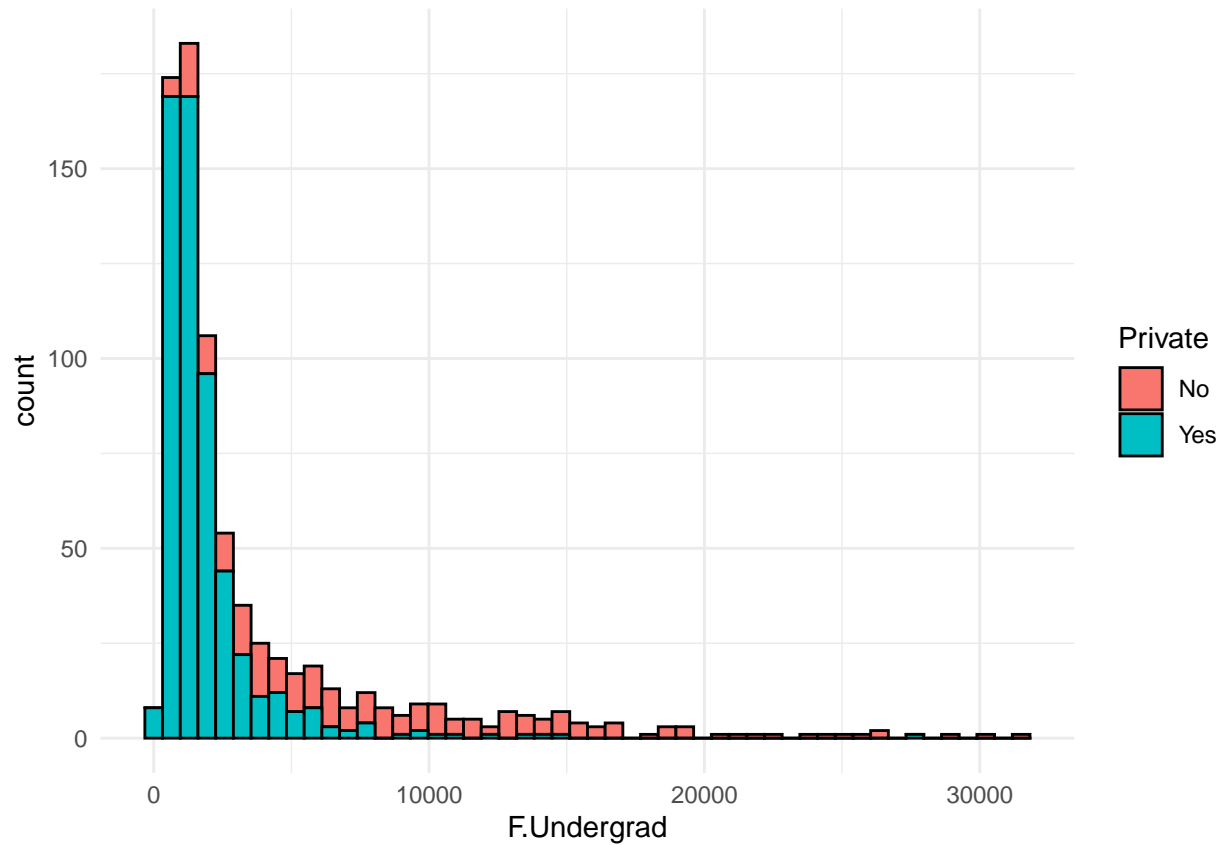
Create a scatterplot of Grad.Rate versus Room.Board, colored by the Private column.

```
ggplot(df, aes(Room.Board, Grad.Rate)) +
  geom_point(aes(color = Private))
```




Create a histogram of full time undergrad students, color by Private.

```
ggplot(df, aes(F.Undergrad, fill = Private)) +  
  geom_histogram(color = "black", bins = 50)
```



Create a histogram of Grad.Rate colored by Private. You should see something odd here.

```
ggplot(df, aes(Grad.Rate, fill = Private)) +  
  geom_histogram(color = "black", bins = 50)
```



What college had a Graduation Rate of above 100% ?

```
rownames(df[df$Grad.Rate > 100, ])
```

```
## [1] "Cazenovia College"
```

Change that college's grad rate to 100%

```
for (i in 1:100){
  if (df$Grad.Rate[i] > 100){
    df$Grad.Rate[i] = 100
  }
}
```

Train Test Split

Split your data into training and testing sets 70/30. Use the caTools library to do this.

```
library(caTools)
```

```
## Warning: package 'caTools' was built under R version 3.6.3
```

```
sample <- sample.split(df$Private, 0.7)

training <- subset(df, sample == T)
testing <- subset(df, sample == F)
```

Decision Tree

Use the rpart library to build a decision tree to predict whether or not a school is Private. Remember to only build your tree off the training data.

```
library(rpart)

treeModel <- rpart(Private ~ ., training, method = "class")
```

Use predict() to predict the Private label on the test data.

```
predict <- predict(treeModel, testing)
```

Check the Head of the predicted values. You should notice that you actually have two columns with the probabilities.

```
head(predict)
```

```
##                No      Yes
## Alverno College    0.02040816 0.97959184
## Andrews University 0.02040816 0.97959184
## Angelo State University 0.94736842 0.05263158
## Antioch University 0.02040816 0.97959184
## Aquinas College    0.02040816 0.97959184
## Arizona State University Main campus 0.94736842 0.05263158
```

Turn these two columns into one column to match the original Yes/No Label for a Private column.

```
predict <- data.frame(predict) %>%
  transmute(Label = ifelse(Yes > No, "Yes", "No"))
```

Now use table() to create a confusion matrix of your tree model.

```
table(as.matrix(predict), testing$Private)
```


```
##
##      No Yes
## No   51  8
## Yes  13 161
```

Use the rpart.plot library and the prp() function to plot out your tree model.

```
library(rpart.plot)
```

```
## Warning: package 'rpart.plot' was built under R version 3.6.3
```

```
prp(treeModel)
```



Random Forest

Now let's build out a random forest model!

Call the randomForest package library

```
library(randomForest)
```

```
## Warning: package 'randomForest' was built under R version 3.6.3
```

Now use `randomForest()` to build out a model to predict Private class. Add `importance=TRUE` as a parameter in the model. (Use `help(randomForest)` to find out what this does.

```
randomForestModel <- randomForest(Private ~ ., training, importance = T)
```

What was your model's confusion matrix on its own training set? Use `model$confusion`.

```
randomForestModel$confusion
```

```
##      No Yes class.error
## No  125  23  0.15540541
## Yes  14 382  0.03535354
```

Grab the feature importance with `model$importance`. Refer to the reading for more info on what `Gini[1]` means.^[2]

```
randomForestModel$importance
```

| ## | | No | Yes | MeanDecreaseAccuracy | MeanDecreaseGini |
|----------------|--------------|--------------|--------------|----------------------|------------------|
| ## Apps | 0.0294866878 | 0.0134061780 | 0.0177890305 | 9.117090 | |
| ## Accept | 0.0370123556 | 0.0137630721 | 0.0200677431 | 14.071934 | |
| ## Enroll | 0.0603772842 | 0.0317175050 | 0.0397756054 | 22.784131 | |
| ## Top10perc | 0.0074196989 | 0.0037018265 | 0.0047190439 | 3.999770 | |
| ## Top25perc | 0.0058087195 | 0.0021203510 | 0.0031284366 | 3.493445 | |
| ## F.Undergrad | 0.1534339586 | 0.0606102382 | 0.0858250636 | 42.745743 | |
| ## P.Undergrad | 0.0391775252 | 0.0065622905 | 0.0153484967 | 16.277205 | |
| ## Outstate | 0.1250765327 | 0.0508757249 | 0.0710933695 | 37.411397 | |
| ## Room.Board | 0.0118341909 | 0.0168489934 | 0.0154349221 | 11.708929 | |
| ## Books | 0.0001775135 | 0.0001670596 | 0.0001739772 | 2.112954 | |
| ## Personal | 0.0047247127 | 0.0002846135 | 0.0014912475 | 3.260299 | |
| ## PhD | 0.0135492166 | 0.0049833572 | 0.0072996940 | 4.692713 | |
| ## Terminal | 0.0063575278 | 0.0054643805 | 0.0057177032 | 4.636212 | |
| ## S.F.Ratio | 0.0347543961 | 0.0099558720 | 0.0167312471 | 18.679402 | |
| ## perc.alumni | 0.0275127440 | 0.0040052434 | 0.0103249880 | 5.889651 | |
| ## Expend | 0.0162117022 | 0.0079925301 | 0.0103153212 | 8.568742 | |
| ## Grad.Rate | 0.0112286512 | 0.0039036740 | 0.0058631821 | 5.231478 | |

Predictions

Now use your random forest model to predict on your test set!

```
predict.rf <- predict(randomForestModel, testing)
```

```
table(predict.rf, testing$Private)
```

```
##
## predict.rf  No Yes
##           No  55  3
##           Yes  9 166
```

It should have performed better than just a single tree, how much better depends on whether you are measuring recall, precision, or accuracy as the most important measure of the model.